

# Parallel Video Mosaic Generation with Metal GPU Acceleration

B12902046 廖昀陽

B12902131 陳柏宇

B13902055 薛閔澤

December 12, 2025

## Abstract

This project implements a high-performance video mosaic generator that reconstructs a live webcam feed using a database of 60,000 tile images (CIFAR-10) in real-time. We compare a multi-threaded CPU implementation using OpenMP against a hardware-accelerated version using Apple Metal. By offloading the computationally intensive feature matching step to the GPU, we achieve a significant speedup, enabling 30+ FPS processing at high resolutions where the CPU implementation struggles at <5 FPS.

## 1 Introduction

Digital photo mosaics are created by replacing grid cells of a target image with smaller "tile" images that best match the color and structure of the original region. Applying this to video in real-time presents a massive computational challenge. For a grid of  $100 \times 75$  cells (7,500 total) and a database of 60,000 tiles, a naïve implementation requires performing  $7,500 \times 60,000 = 450$  million comparisons per frame. At 30 FPS, this equates to 13.5 billion comparisons per second.

This project explores parallel optimization techniques to solve this problem, progressively moving from a multi-threaded CPU approach to a massively parallel GPU solution on Apple Silicon (M3).

## 2 Methodology

### 2.1 Evolution of Approaches

We explored three distinct algorithmic approaches to solve the mosaic generation problem:

### 2.1.1 1. Naive Greedy Approach ('video\_mosaic.cpp')

The simplest implementation iterates through each grid cell and selects the tile with the minimum Euclidean distance in RGB space.

1. **Pros:** Fast ( $O(N \cdot M)$  where  $N$  is cells,  $M$  is tiles).
2. **Cons:** Result is visually flat; ignores structure/edges. Tiles are repeated frequently, making the mosaic look like a low-resolution sampling rather than a collage.

### 2.1.2 2. Global Optimization ('video\_mosaic\_optimal.cpp')

To improve visual quality and reduce repetition, we implemented a cost-matrix based approach (approximating the Hungarian Algorithm). We compute a cost matrix for all cell-tile pairs and attempt to minimize the total cost while enforcing unique tile usage.

1. **Pros:** High artistic quality; no varying tile repetition.
2. **Cons:** Extremely computationally expensive. Solving the assignment problem for 7,500 cells and 60,000 tiles is infeasible for real-time video (seconds per frame), even with greedy approximations.

### 2.1.3 3. Edge-Aware Matching (Final Approach)

To balance performance and visual fidelity, we adopted the "Edge-Aware" metric. It allows tile repetition (unlike the optimal approach) but enforces structural consistency (unlike the naive approach).

- **Metric:**  $Cost = w_c \cdot Dist_{color} + w_e \cdot Dist_{edge}$ .
- **Edge Features:** We use Sobel operators to extract gradient magnitude and direction. A tile only matches a grid cell if they share similar edge orientations, preserving the shapes of objects in the video.

## 2.2 CPU Optimization (OpenMP)

Our optimized CPU implementation ('video\_mosaic\_edge\_optimized') uses two key strategies:

- **Global Feature Extraction:** Instead of resizing and computing Sobel gradients for each of the 7,500 grid cells individually (which involves redundant computation at boundaries), we resize the full input frame once and compute global gradient maps. Grid features are then sliced directly from these maps.
- **Parallel Search:** We use OpenMP ('#pragma omp parallel for') to parallelize the search. Each thread handles a subset of grid cells.

## 2.3 GPU Acceleration (Metal)

To overcome CPU compute limits, we implemented a custom Metal compute kernel.

- **Wait-Free Parallelism:** The problem is "embarrassingly parallel". Each grid cell can find its best matching tile independently of all others.
- **Compute Kernel:** We launch one GPU thread per grid cell. Each thread iterates through the entire 60,000-tile database stored in GPU private memory, finds the tile index with the minimum cost, and writes it to an output buffer.
- **Unified Memory:** On Apple Silicon, we leverage the unified memory architecture. The CPU writes grid features to a buffer, and the GPU reads them directly, minimizing data transfer overhead compared to discrete GPUs.
- **Memory Layout:** We use a Structure-of-Arrays (SoA) layout padded to ‘float4’ (16 bytes) alignment to ensure coalesced memory access and vectorization on the GPU.

## 3 Experimental Results

### 3.1 Hardware Setup

- **System:** Apple MacBook Pro (M3 Chip)
- **Memory:** Unified Memory Architecture
- **Dataset:** CIFAR-10 (60,000 images, 32x32 pixels)

« « < HEAD

## 4 Experimental Results

### 4.1 Hardware and Workload Setup

- **System:** Apple MacBook Pro (M3 Chip with Unified Memory).
- **Dataset:** CIFAR-10 (60,000 images,  $32 \times 32$  pixels).
- **Workload Configurations:**
  - **Medium Grid:**  $60 \times 45$  cells (2,700 total). Requires  $2,700 \times 60,000 = 162$  million comparisons per frame.
  - **Ultra Grid:**  $100 \times 75$  cells (7,500 total). Requires  $7,500 \times 60,000 = 450$  million comparisons per frame.

## 4.2 Performance Benchmark

We evaluated the frame rate (FPS) of three implementations:

1. **CPU 1T**: Optimized CPU code running on 1 thread.
2. **CPU 8T**: Optimized CPU code running on 8 threads (OpenMP).
3. **Metal**: GPU accelerated implementation.

Configuration	Workload (Grid Size)	FPS	Speedup
CPU 1T	Medium ( $60 \times 45$ )	0.73	1.0x (Baseline)
CPU 4T	Medium ( $60 \times 45$ )	2.22	3.0x
CPU 8T	Medium ( $60 \times 45$ )	2.41	3.3x
<b>Metal GPU</b>	<b>Medium (<math>60 \times 45</math>)</b>	<b>10.27</b>	<b>14.1x</b>
CPU 8T	Ultra ( $100 \times 75$ )	0.86	—
<b>Metal GPU</b>	<b>Ultra (<math>100 \times 75</math>)</b>	<b>7.09</b>	<b>8.2x (vs CPU 8T)</b>

Table 1: Performance comparison. The Metal GPU implementation maintains real-time performance even at Ultra resolution where the CPU implementation drops to <1 FPS.

## 4.3 Analysis

The CPU implementation shows limited scalability, peaking at 2.41 FPS on Medium settings with 8 threads. The massive computational demand of performing 450 million metric evaluations per frame (Ultra settings) brings the CPU to a crawl (0.86 FPS).

The Metal implementation drastically outperforms the CPU, achieving **10.27 FPS** on Medium settings (14x speedup over single-threaded CPU). Even on the demanding Ultra setting, it maintains **7.09 FPS**, providing a smooth visual experience that is impossible with the CPU-only approach. The slight drop in GPU performance at Ultra settings primarily comes from the overhead of rendering and capturing larger video frames, rather than the tile matching kernel itself. =====

## 4.4 Performance Benchmark

We compared the implementation on two grid sizes: Medium ( $60 \times 45$ ) and Ultra ( $100 \times 75$ ).

<b>Configuration</b>	<b>Grid Size</b>	<b>FPS</b>	<b>Speedup</b>
CPU Optimized (1 Thread)	Medium	0.71	1.0x (Baseline)
CPU Optimized (4 Threads)	Medium	2.14	3.0x
CPU Optimized (8 Threads)	Medium	2.46	3.5x
<b>Metal GPU</b>	<b>Medium</b>	<b>11.87</b>	<b>16.7x</b>
CPU Optimized (8 Threads)	Ultra	0.86	—
<b>Metal GPU</b>	<b>Ultra</b>	<b>10.86</b>	<b>12.6x (vs CPU 8T)</b>

Table 2: Performance comparison between CPU and GPU implementations. Speedup is calculated relative to the single-threaded CPU baseline for Medium, and relative to the best CPU result for Ultra.

#### 4.5 Analysis

The CPU implementation scales linearly up to 4 threads, but sees diminishing returns at 8 threads (2.46 FPS), likely due to memory bandwidth saturation from the constant reloading of tile data into cache.

The Metal implementation yields a dramatic speedup. Notably, the performance drop from "Medium" (11.87 FPS) to "Ultra" (10.86 FPS) is very small on the GPU, whereas the CPU struggles significantly (dropping to <1 FPS). This indicates that the GPU implementation is not compute-bound by the matching logic, but rather limited by the serial overhead of video capture and rendering, proving it effectively solved the computational bottleneck. »»>  
33c260b967bd071c5d0d1a143e3b5242df9fc7e

### 5 Conclusion

We successfully implemented a real-time edge-aware video mosaic generator. While CPU optimizations provided a 4.5x speedup over single-threaded execution, they fell short of real-time performance for high-resolution grids. The Metal GPU implementation solved this bottleneck, enabling smooth >30 FPS processing even at "Ultra" grid settings, demonstrating the power of hardware acceleration for massive pattern matching tasks.